

SEALED MEMO: Technical Agent

Candidate: rohan_malhotra | Created: 2026-08-28 09:33:11 UTC

Score (1-10)	8
Confidence	HIGH

Verdict Summary

Direct architectural match for Cargonet's stack. Hands-on experience designing planner/executor/reviewer multi-agent pipelines for freight exception handling, prompt engineering, and SLM/GPT-4 cost routing.

Identified Strengths (3)

- [R-EXP-01] Designed planner/executor/reviewer multi-agent exception handling engine handling freight exceptions
- [R-EXP-02] Implemented cost-effective model routing across GPT-4 and open-weight SLMs reducing inference costs ~30%
- [R-EXP-05] Experience with LangChain and Pinecone vector search over carrier rate documents

Identified Gaps / Concerns (2)

- [T-A3] Vague on reviewer agent evaluation metrics, relying only on untracked override rates
- [T-A4] Model routing was tuned heuristically as things broke rather than via rigorous evaluation benchmarks