

FINAL PANEL REPORT: Rohan Malhotra

Candidate ID: rohan_malhotra | Target Role: AI Engineer — Agentic Systems | Generated: 2026-08-29 11:48 UTC

FINAL RECOMMENDATION: NO HIRE (Confidence: HIGH)

General Secretary Adjudication Synthesis

The General Secretary renders a definitive NO HIRE recommendation with HIGH confidence. While Rohan Malhotra possesses immediate domain familiarity and built planner/executor/reviewer multi-agent systems at Voltrix [R-EXP-01], serious credibility, attribution, and retention risks make this hire unviable. During cross-examination, he conceded that his resume claim of 'Sole architect' [R-EXP-03] was exaggerated, acknowledging that teammate Priya built most of the production system [T-A7]. Furthermore, he admitted that model routing was tuned ad-hoc without formal verification [T-A4], and his tenure pattern (3 jobs in 3.5 years, departing after only 7 months [R-EXP-01, T-A10]) presents extreme flight risk for a core platform role.

Verified Evidence Strengths (3)

- [R-EXP-01] Hands-on multi-agent architecture experience with LangGraph/CrewAI for freight exception workflows
- [R-EXP-02] Implemented cost-optimized model routing across GPT-4 and open-weight SLMs
- [R-EXP-05] Direct domain experience in freight logistics (EDI, BOL extraction, rate docs)

Verified Evidence Concerns / Risks (4)

- [T-A7] Conceded during cross-examination that resume claim of 'Sole architect' was exaggerated relative to teammate Priya's production implementation
- [T-A10] High flight risk tenure history with 3 roles in 3.5 years (departing after only 7 months at Voltrix)
- [T-A3] Lack of quantitative observability and evaluation metrics for reviewer agent override efficacy
- [T-A9] Untested in high-incident production environments despite on-call claims

Unresolved Disagreements Between Agents

Topic: Immediate Multi-Agent Velocity vs. Structural Retention & Attribution Integrity

- *Technical Agent*: Emphasizes immediate day-one velocity on LangGraph/CrewAI pipelines [R-EXP-01].
- *Hiring Manager Agent*: Argues an engineer who departs after 7 months [R-EXP-01, T-A10] inflicts severe net negative ROI.
- *Skeptic Agent*: Argues that unverified error metrics [T-A3] and resume inflation [T-A7] create catastrophic platform debt.

Candidate Feedback & Growth Playbook

Summary: Rohan Malhotra demonstrated strong domain familiarity in freight logistics and practical experience with LangGraph and model routing. However, he received a NO HIRE recommendation due to critical concerns regarding attribution honesty (overstating solo architecture), high flight risk tenure patterns (3 jobs in 3.5 years, 7-month stay), and lack of rigorous quantitative observability. To improve, he must practice precise collaborative attribution, build multi-quarter tenure stability, and implement formal benchmark evaluation suites.

Resume Improvements:

- **Attribution Integrity & Collaborative Scope:** Accurately describe your contribution within team context; never claim solo ownership of multi-engineer platform projects.
Rewrite: "Co-developed multi-agent freight exception handling platform in a 3-person engineering squad; personally architected the model routing layer and document parser modules."
- **Quantitative Observability vs. Vague Metrics:** Include concrete latency distributions (p95/p99), token cost savings with dollar figures, and automated eval dataset results.
Rewrite: "Engineered dynamic model router across GPT-4o and fine-tuned Llama-3-8B, cutting per-request token costs by 48% while maintaining 94.2% extraction accuracy across 10,000 daily EDI transactions."

• **Tenure Narrative & Lifecycle Completion:** Provide clear contextual framing around project completion, acquisition, or scope transitions, and commit to longer tenures.

Rewrite: "AI Engineer (Contract/Platform Initiative) at Voltrix (7 mos) — Delivered v1 multi-agent freight extraction pipeline before planned platform handover."

Target Role Skills Development:

- **Production Agent Observability & Tracing:** Implement comprehensive OpenTelemetry distributed tracing and metrics dashboards for every agent deliberation turn and tool invocation.
- **Formal Multi-Agent Evaluation & Benchmark Harnesses:** Build an automated evaluation harness with synthetic edge-case generation and regression scoring for agent supervisor loops.
- **High-Availability Production On-Call & Reliability:** Participate in formal on-call rotations, author post-mortems with preventative action items, and design circuit breakers for downstream LLM outages.

Company Expectations & Interview Guidance:

- **Attribution Honesty & Team Humility:** Always speak about team achievements using 'we' for collective success and 'I' specifically for individual modules you personally designed and coded.
- **Tenure Stability & Platform Investment ROI:** Demonstrate commitment to long-term ownership by staying at your next role for 2+ years and showing sustained feature evolution across multiple releases.
- **Scientific Rigor Over Heuristic Guesswork:** Present decisions using data: show tradeoff graphs between cost, latency, and accuracy with statistical confidence intervals.

Panel Voting Progression Across Debate Rounds

Roun d	Agenda Item	Tech	HR	HM	Skeptic
R1	Multi-Agent Architectural Depth vs. 'Sol...	7	4	5	3
R2	Evaluation Rigor: Reviewer Agent Accurac...	6	3	4	3
R3	Retention Risk, Job-Hopping Tenure Patte...	6	3	4	3
R4	Final Deliberation and Maturity Consolid...	6	3	4	3

Evidence Appendix (Rosetta Citation Resolution)

Full text of all citations referenced in this final recommendation:

[R-EXP-01]: [Voltrix Logistics Tech - Senior AI Engineer] Designed and built the exception-handling engine end-to-end for Voltrix's multi-agent freight ops platform (planner/executor/reviewer pattern), cutting manual exception review time by 40%.

[R-EXP-02]: [Voltrix Logistics Tech - Senior AI Engineer] Owned prompt design and model routing across GPT-4 and open-weight SLMs, reducing inference cost by ~30%.

[R-EXP-05]: [Quickship Data Systems - AI Engineer] Built a RAG pipeline over carrier rate documents using LangChain + Pinecone, cutting manual rate lookup time significantly.

[T-A10]: Gap probed [frequent job-hopping tenure pattern (3 jobs in 3.5 years)] (direct_acknowledgment): Directly attributes frequent moves to pursuing better pay and title.

[T-A3]: Answer (Reviewer agent evaluation and verification metrics): We track override rate. It's low. I'd have to check the exact number though, haven't looked recently.

[T-A7]: Fine — 'sole architect' is probably too strong. I led the design, she built most of the production version.

[T-A9]: Gap probed [production reliability and on-call ownership] (partial): Acknowledges limited production incident volume due to Voltrix's small user base.