

FINAL PANEL REPORT: Rohan Malhotra

Candidate ID: rohan_malhotra | Target Role: AI Engineer — Agentic Systems | Generated: 2026-08-28 06:57 UTC

FINAL RECOMMENDATION:	NO_HIRE	CONFIDENCE LEVEL:	HIGH
-----------------------	---------	-------------------	------

General Secretary Adjudication Synthesis

The General Secretary renders a definitive NO-HIRE recommendation with HIGH confidence. While Rohan Malhotra possesses relevant conceptual familiarity with planner/executor/reviewer architectures [R-EXP-01], serious evidentiary defects disqualify his candidacy. Under cross-examination in [T-A7], he walked back his resume claim of 'sole architect' [R-EXP-03], conceding that teammate Priya implemented most of the production code. Furthermore, he demonstrated a complete absence of evaluation rigor, failing to track reviewer override rates [T-A3] and tuning model routing heuristically as things broke [T-A4]. Most critically, having held 3 jobs in 3.5 years [T-A10] driven purely by short-term compensation hopping, hiring him represents an untenable flight risk and negative net ROI.

Override Motion Record (PRD §10)

Motion filed by Technical Agent: "Candidate has direct, immediate architectural familiarity with planner/executor/reviewer freight pipelines and can ship features on day one [R-EXP-01]."

Outcome: 1/4 votes in favor (FAILED - Supermajority not reached).

Rationale: Motion failed (1/4 votes in favor). Panel overwhelmingly concluded that day-one speed cannot compensate for material integrity concerns, absent evaluation metrics, and severe 6-month flight risk.

Verified Evidence Strengths (3)

- [R-EXP-01] Hands-on familiarity designing planner/executor/reviewer exception handling patterns for freight ops
- [R-EXP-02] Implemented cost-based model routing across GPT-4 and open-weight SLMs reducing inference expense
- [R-EXP-05] Experience building RAG pipelines over carrier rate documents with LangChain and Pinecone

Verified Evidence Concerns / Risks (6)

- [T-A7] Material resume misrepresentation: claimed 'sole architect' on resume but admitted in interview that teammate Priya built most of production code
- [T-A10] Severe retention flight risk with 3 jobs in 3.5 years, explicitly motivated by short-term title and salary hops
- [T-A3] Zero evaluation rigor: unable to provide metrics or override rates for production reviewer agent
- [T-A4] Model routing was tuned heuristically as things broke without formal evaluation sets or regression benchmarks
- [T-A6] Defensive responses and friction regarding credit sharing on engineering projects
- [T-A9] Dismissive of production on-call operational rigor due to small past user bases

Unresolved Disagreements Between Agents

- Topic: Day-One Multi-Agent Domain Velocity vs. Flight Risk & Integrity Deficit**
- *Technical Agent:* Emphasizes immediate productivity on planner/executor/reviewer freight workflows [R-EXP-01].
 - *Hiring Manager Agent:* Argues an engineer who departs after 7 months [R-EXP-01, T-A10] inflicts severe net negative ROI.
 - *Skeptic Agent:* Argues that unverified error metrics [T-A3] and resume inflation [T-A7] create catastrophic platform debt.

Panel Voting Progression Across Debate Rounds

Round	Agenda Item	Tech	HR	HM	Skeptic
R1	Multi-Agent Architectural Depth vs. 'Sol...	7	4	5	3

R2	Evaluation Rigor: Reviewer Agent Accurac...	6	3	4	3
R3	Retention Risk, Job-Hopping Tenure Patte...	6	3	4	3
R4	Final Deliberation and Maturity Consolid...	6	3	4	3

Evidence Appendix (Rosetta Citation Resolution)

Full text of all citations referenced in this final recommendation:

[R-EXP-01]: [Voltrix Logistics Tech - Senior AI Engineer] Designed and built the exception-handling engine end-to-end for Voltrix's multi-agent freight ops platform (planner/executor/reviewer pattern), cutting manual exception review time by 40%.

[R-EXP-02]: [Voltrix Logistics Tech - Senior AI Engineer] Owned prompt design and model routing across GPT-4 and open-weight SLMs, reducing inference cost by ~30%.

[R-EXP-05]: [Quickship Data Systems - AI Engineer] Built a RAG pipeline over carrier rate documents using LangChain + Pinecone, cutting manual rate lookup time significantly.

[T-A10]: Gap probed [frequent job-hopping tenure pattern (3 jobs in 3.5 years)] (direct_acknowledgment): Directly attributes frequent moves to pursuing better pay and title.

[T-A3]: Answer (Reviewer agent evaluation and verification metrics): We track override rate. It's low. I'd have to check the exact number though, haven't looked recently.

[T-A4]: Answer (Model routing and cost optimization approach): Cost-based. Simple stuff to the SLM, harder reasoning to GPT-4. No formal study, just tuned it as things broke.

[T-A6]: I designed it. Priya did a lot of the implementation, I reviewed her PRs. I was the architect.

[T-A7]: Fine — 'sole architect' is probably too strong. I led the design, she built most of the production version.

[T-A9]: Gap probed [production reliability and on-call ownership] (partial): Acknowledges limited production incident volume due to Voltrix's small user base.